

Introduction to Large Language Models(LLM)

Outline

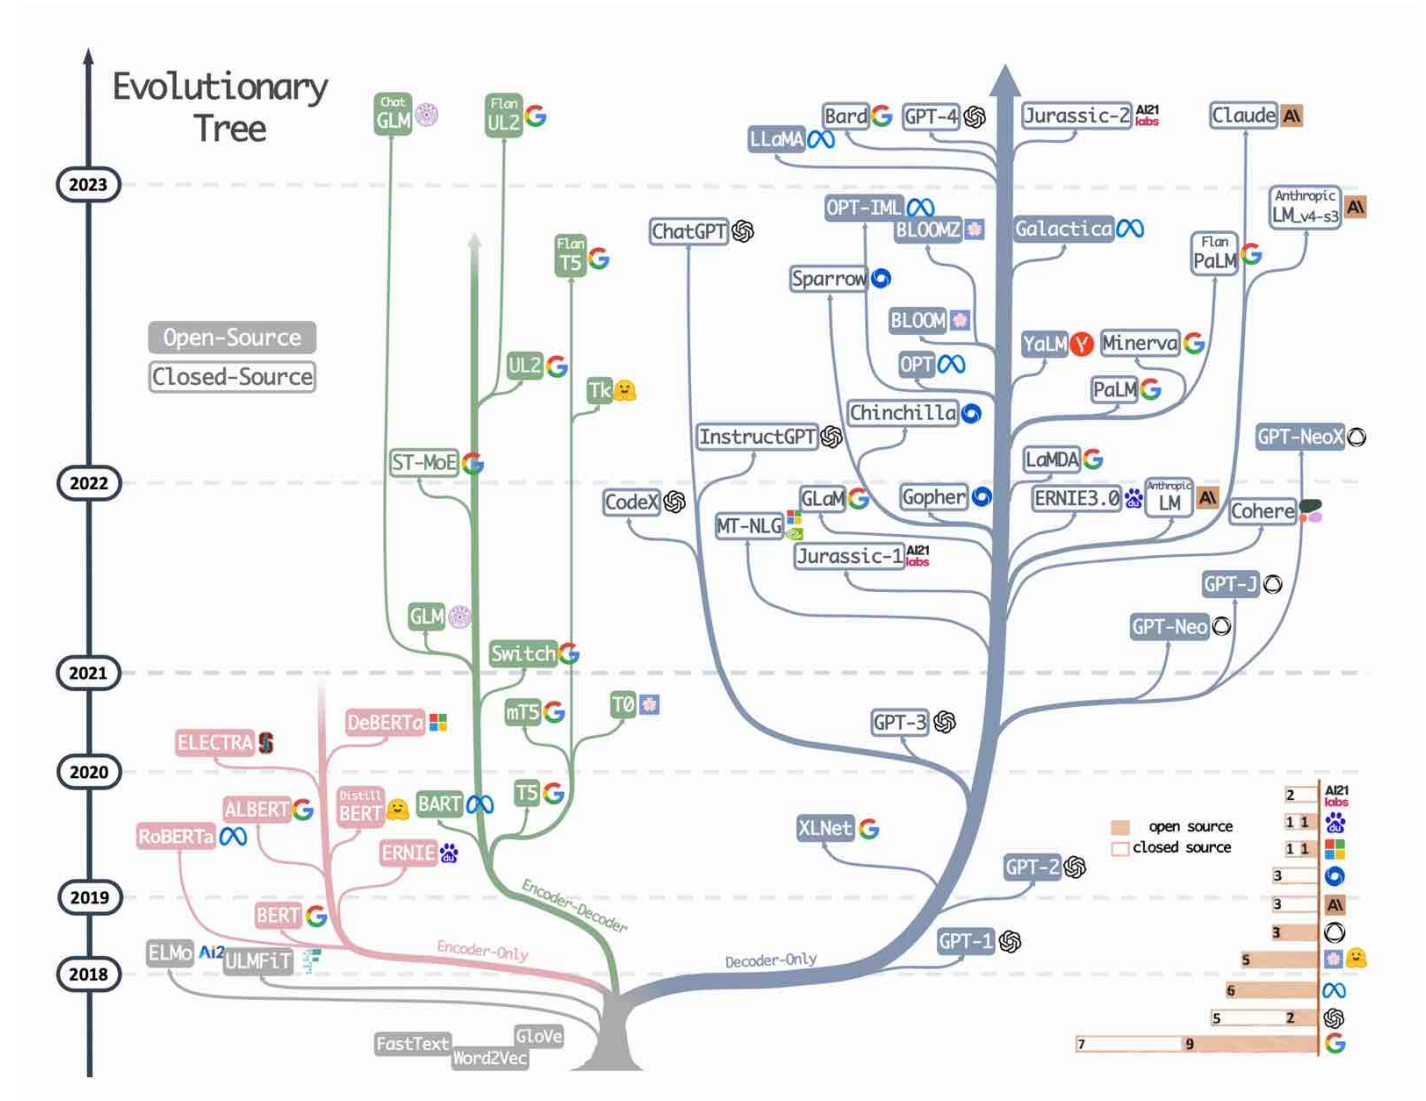
1. Overview from 30,000 feet above
2. Transformer in nutshell
3. Pretraining - Parallel paradigm
4. Finetuning - Parameter efficient finetuning
5. Steering the decoding of LLM - Prompting
6. Augmentation and Plugins

Overview from 30000 feet above

- Paradigm transition in AI
 - From: **training**(specific) -> **prediction**(specific)
 - To: **pretraining**(general) -> **finetuning**(general/specific) -> **in-context prompt**(specific)
- Primary steps of new paradigm
 - Pretraining with self-supervised learning
 - Finetuning on instruction from multiple domains
 - Application by steering/prompt the decoding process of LLM
- Where are we to AGI?
 - From explanation to prediction
 - From correlation to causality

A LLMs Evolution Tree

- Decoder Only
- Encoder Only
- Encoder-Decoder

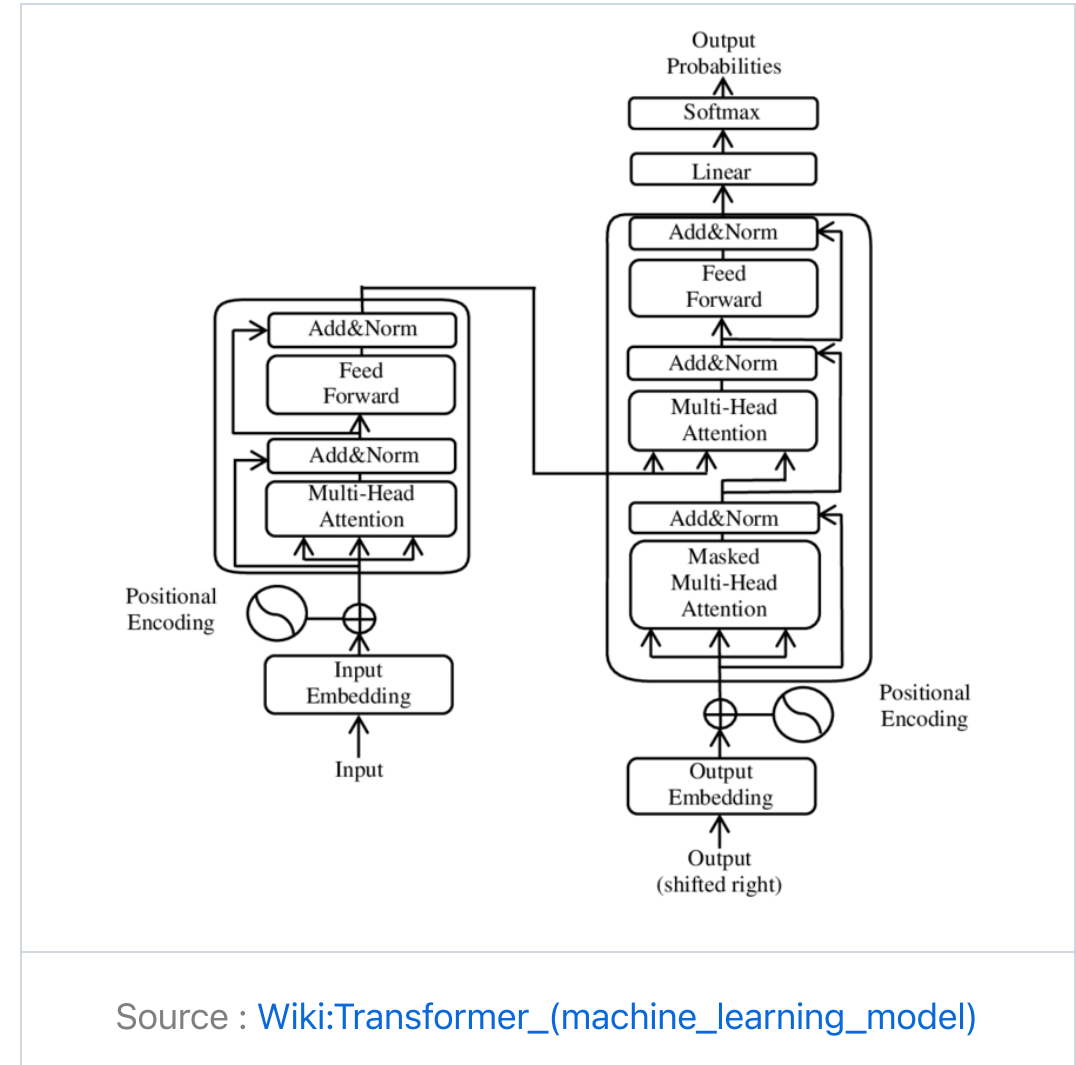


Transformer in nutshell

- Transformer modules
- Aspects of alternative
- Parameter concentration
- Computation concentration

Transformer modules

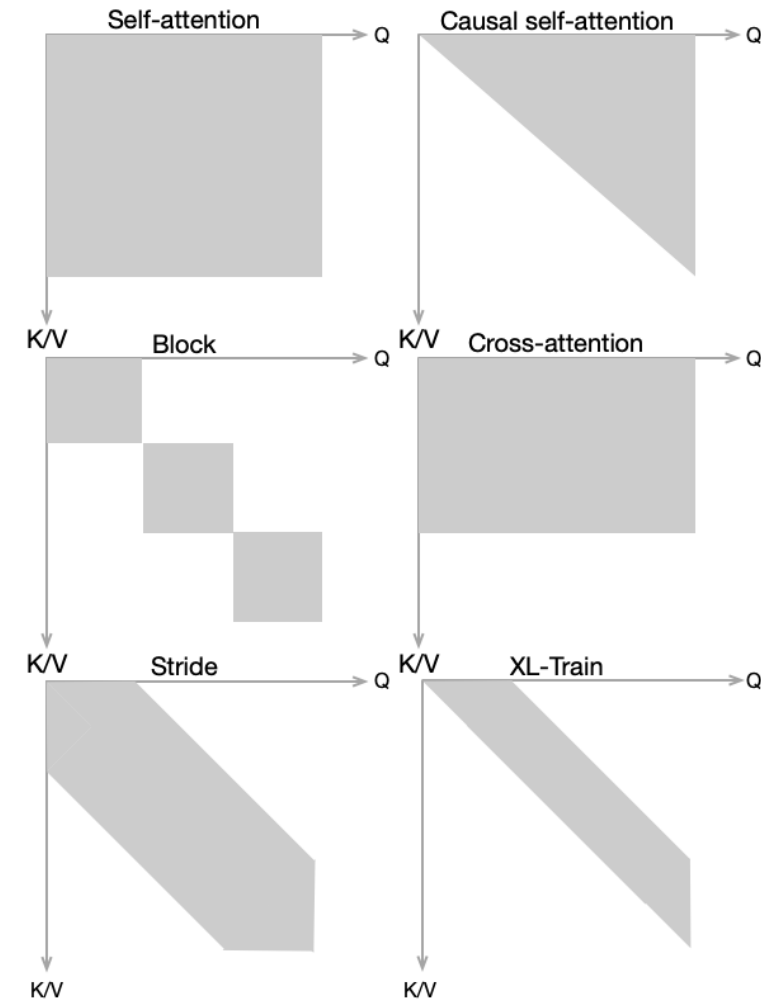
- Token & positional embedding :
 - $E \in R^{V \times d}, P \in R^{T \times d}$
- Multi-head attention
 - Self-attention & Cross-attention
 - Weight matrix:
 $W_Q, W_K, W_V \in R^{d \times d_h}$
 - Head projection:
 $W_O \in R^{d \times d}$
- ResNet & LayerNorm
- Feedforward Network
 - $W_1 \in R^{d \times 4d}, W_2 \in R^{4d \times d}$
- Output head: task related



Source : [Wiki:Transformer_\(machine_learning_model\)](https://en.wikipedia.org/wiki/Transformer_(machine_learning_model))

Aspects of alternative

- Efficient Transformer (for long sequence)
 - Coarse sequence resolution:
 - Block/Stride/Clustering/Neural Memory
 - TransformerXL
 - Attention matrix approximation:
 - Linear Transformer
- LLMs specific
 - **Encoder vs Decoder vs Encoder-Decoder**
 - Pretraining objective
 - Positional encoding
 - Input or output LayerNorm
 - Activation



Parameter concentration

Decoder only Transformer(GPT)

- Parameter size:
 - Embedding: $(V + T) \times d$
 - Attention: $L \times (3 \times d \times d_h \times H + (d_h \times H) \times d) = 4Ld^2$
 - FFN: $L \times ((d \times 4d + 4d) + (4d \times d + d)) \approx 8Ld^2$
- On GPT3-175B:

Total	PE	TE	Attn	FFN
174,597M	25M(0.01%)	617M(0.35%)	57,982M(33.21%)	115,970M(66.42%)

Computation concentration

Decoder only Transformer

Per-token calculation:

- QKV+project: $2 \times L \times (3 \times H \times h_d \times d + (H \times h_d) \times d) = 2 \times 4Ld^2$
- Attention: $2 \times L \times T \times d$
- FFN: $2 \times L \times ((d \times 4d + 4d) + (4d \times d + d)) \approx 2 \times 8Ld^2$

Model training flops utilization(MFU):

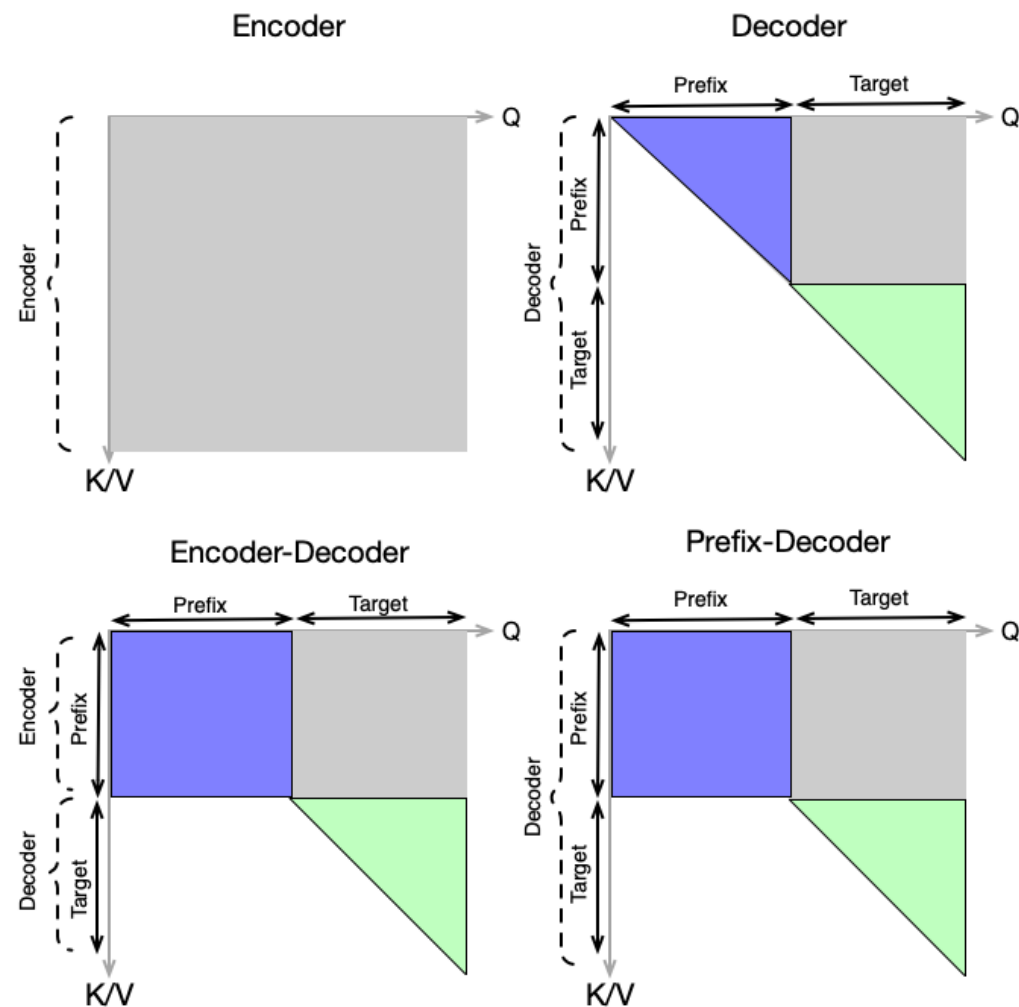
- Forward and backward: $(1 + 2) \times (2N + 2LTd) \approx 6N$, where $N \approx 12Ld^2$
- Theoretical peak throughput: $\frac{P}{6N+6LTd}$
- $\text{MFU} = \frac{\text{Observed throughput}}{\text{Theoretical peak throughput}}$

Pretraining

- Training objectives
- Text Corpus
- Parallel strategies
- Results & Evaluation

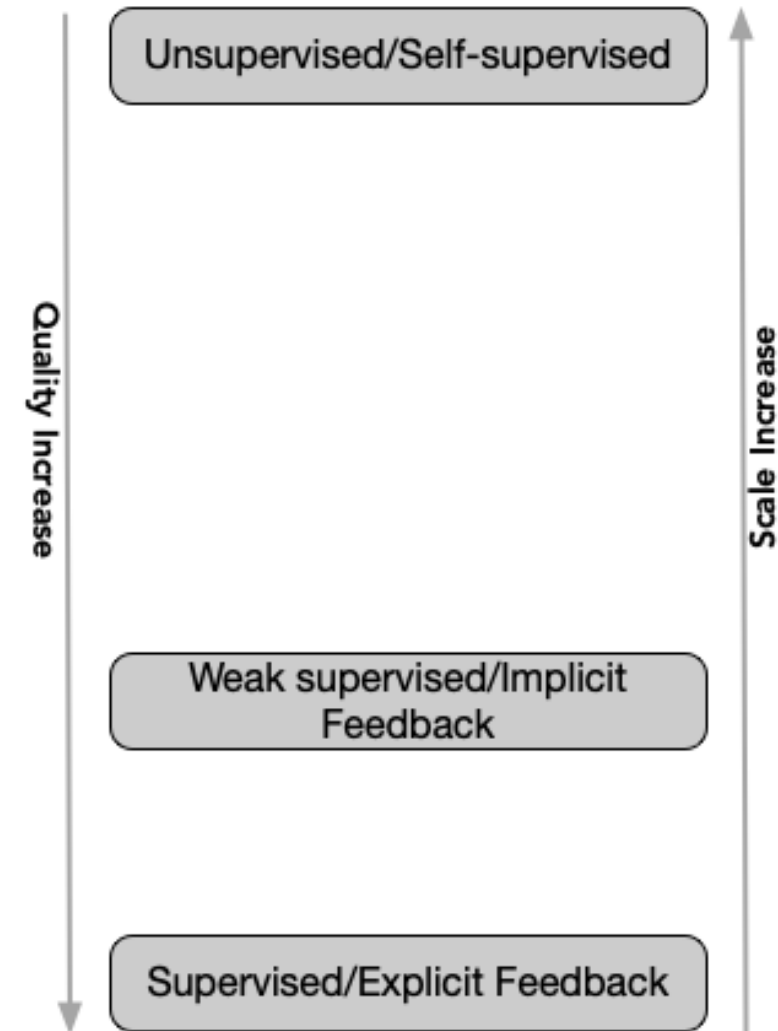
Training architecture & objectives

- Architecture
 - Encoder: BERT series
 - Encoder-Decoder: T5(11B)
 - (Causal-)Decoder: GPTs/Ernie3.0 Titan
 - Prefix-Decoder: GLM
- Objectives
 - Masked LM: BERT/GLM/T5
 - Auto-regressive LM: T5/GPTs/Ernie3.0 Titan
 - Multi-task pretraining: GLM/Ernie3.0 Titan



Text Corpus

- Unsupervised text
 - [BookCorpus](#) : 11,000; [Gutenberg](#) : 70,000
 - [OpenWebText](#): 8M outlinks of Reddit.com
 - [Common Crawl](#); [C4](#)
 - Code: [BigQuery Github](#)
- Weak supervised text
 - [Reddit TL;DR](#); [PushShift.io](#) [Reddit](#): Posts
 - [StackExchange](#) : Question & Answer w/ score
- Supervised text: task related
 - ~16 NLP tasks related datasets (Sentiment/QA/Reasoning, etc.)
 - Human answer to prompt: InstructGPT



Baidu Ernie 3.0 Titan: 260B

- 2 Types of transformer modules
 - Universal module
 - Task specific module
- 3 Levels of pretraining tasks
 - Word aware: Knowledge integrated LM
 - Structure aware: sentence reordering task
 - Knowledge aware: controllable LM task
- 4D hybrid parallelism for training
 - Shared data parallel with ZeRO(2D)
 - Intra-layer tensor parallel(D)
 - Inter-layer pipeline parallel(D)

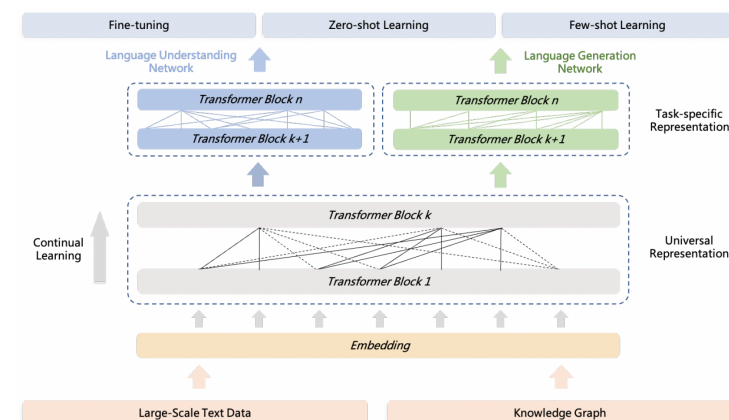


Figure 1: The framework of ERNIE 3.0.


Parallel strategies

- Why bother?
- Three parallel paradigms
 - Data parallel: ZeRO
 - Model parallel
 - Tensor parallel: Megatron-LM
 - Pipeline parallel: GPipe
- Combined implementations: DeepSpeed/ColossalAI

Why bother?

- Too big to fit in single GPU memory
 - 175B: $\sim (2 + 2 + 3 \times 4) \times 175 = 2800\text{GB}$ for mixed-precision training
 - Parameter & gradient(FP16): parameter(W), gradient($g, \frac{\partial L}{\partial W}$)
 - Optimizer State(FP32): parameter, momentum(m), variance(v)
 - Activation(FP16): $2 \times (1 + 4 + 1) \times d \times B \times T \times L$
 - A100 Spec:
 - GPU memory: 80GB
 - GPU memory bandwidth: 2039GB/s; NVLink: 600GB/s; PCIe 4.0: 64GB/s
 - TF32: 156TFlops
- Speedup
 - Scales linearly with # of GPU cores?

Basics on Forward & Backward and Parallel Ops



Forward:

$$h_0 = \sigma(z_0)$$

$$z_1 = W^T h_0 + b$$

$$h_1 = \sigma(z_1)$$

Backward:


$$\left(\frac{\partial L}{\partial h_1}, W\right) \rightarrow \frac{\partial L}{\partial h_0}$$

$$\left(\frac{\partial L}{\partial h_1}, h_0\right) \rightarrow \Delta \frac{\partial L}{\partial W}$$

$$m \leftarrow \beta_1 m + (1 - \beta_1) \frac{\partial L}{\partial W}$$

$$v \leftarrow \beta_2 v + (1 - \beta_2) \left(\frac{\partial L}{\partial W}\right)^2$$

$$W \leftarrow W - \frac{\alpha}{\sqrt{\hat{v} + \epsilon}} \hat{m}$$



Data parallel: from DDP to FSDP(ZeRO)

- *Pesudo code* for DDP and FSDP

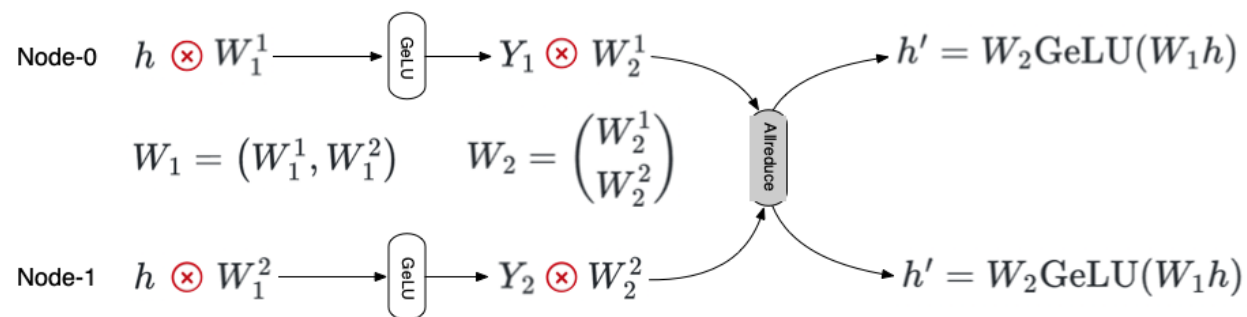
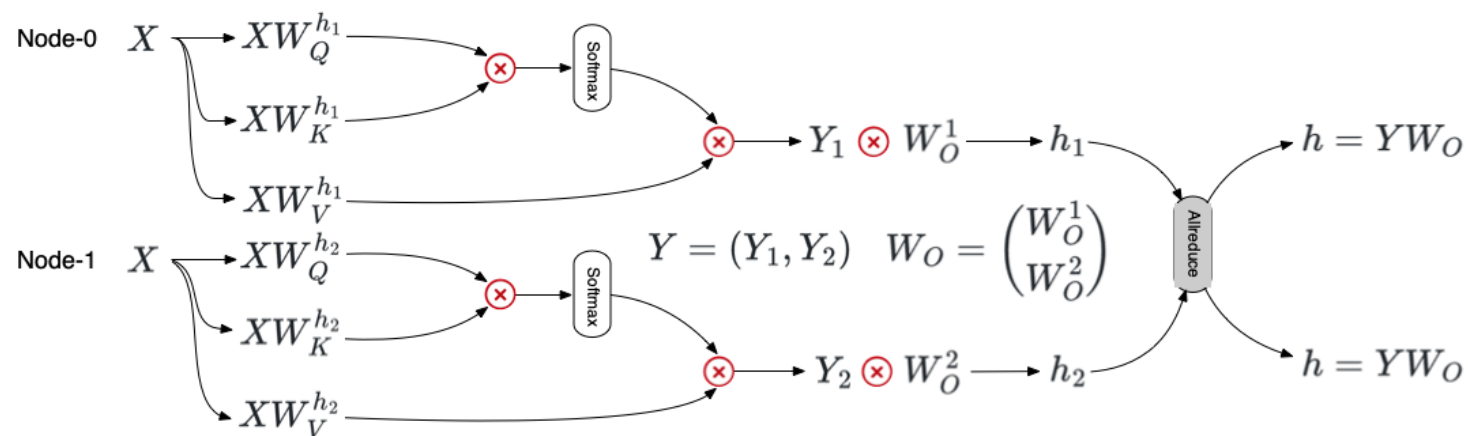
```
# forward pass :  
for layer_i in layers:  
    forward pass for layer_i  
# backward pass :  
for layer_i in layers:  
    backward pass for layer_i  
    full: all-reduce gradients for layer_i  
    full: update momentum & variance  
    full: update weights
```

```
# forward pass :  
for layer_i in layers:  
    all-gather full weights for layer_i  
    forward pass for layer_i  
    discard full weights for layer_i  
# backward pass:  
for layer_i in layers:  
    all-gather full weights for layer_i  
    backward pass for layer_i  
    discard full weights for layer_i  
    part: reduce-scatter gradients for layer_i  
    part: update momentum & variance  
    part: update weights
```

- Advantages of ZeRO
 - Parameter & gradient and optimizer states evenly shard to N nodes
 - Computation and communication overlaps

Tensor parallel: Megatron-LM

- W_Q, W_K, W_V partition by head(col)
- W_O partition by row
- W_1 partition by col, W_2 by row
- Backward Allreduce for gradient

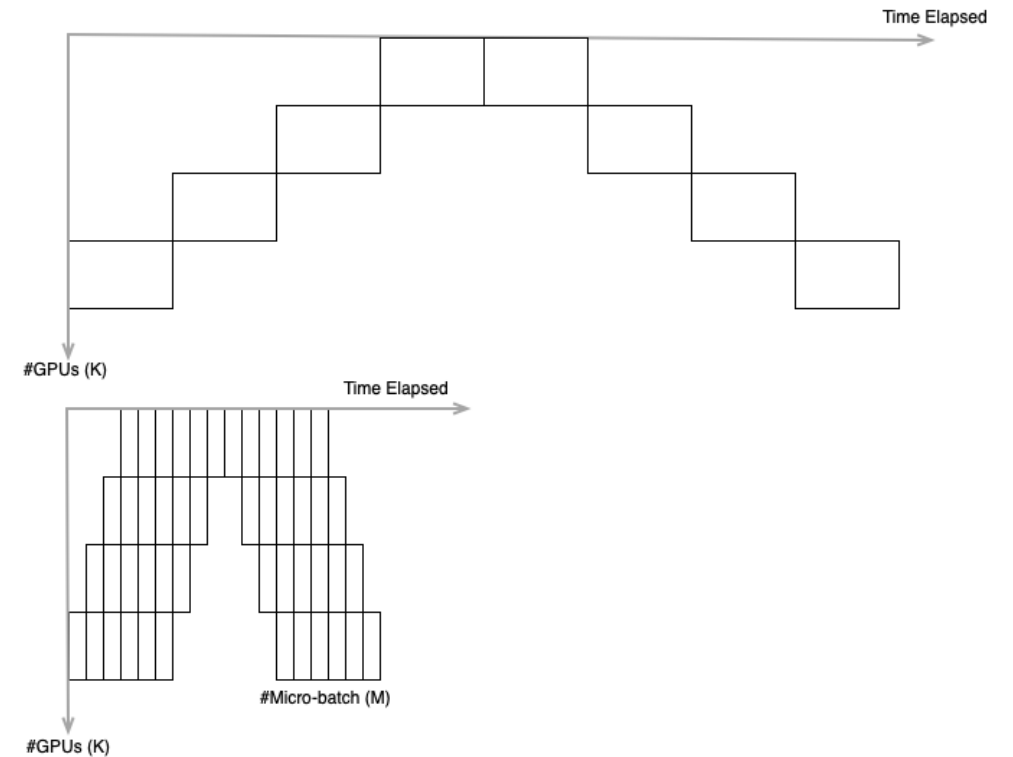



Performance comparison of ZeRO and Megatron-LM

- Experiment hardware:
 - Megatron-LM: 32 DGX-2H servers: 512 V100, 32GB GPUs
 - ZeRO: 25 DGX-2 servers: 400 V100, 32GB GPUs
- TFlops Results:
 - Megatron-LM: 15.1 PFlops
 - 76% scaling efficiency for single GPU 39TFlops(30% of peak Flops)
 - ZeRO: 15 PFlops, with fewer GPU cores

Pipeline parallel: GPipe

- Layer-wise model partition
- Re-materialization: output activation stored and communicated
- Pipeline reduce bubble ratio from $\frac{K-1}{K}$ to $\frac{K-1}{M+K-1}$



Combined implementations: Megatron-DeepSpeed/ColossalAI

Megatron-DeepSpeed

- Data/Model parallel supported(3D)
- ZeRO-Offload
- Sparse attention
- 1-bit Adam and 0/1 Adam
- MoE specific parallelism
- RLHF Demo: deepspeed-chat

ColossalAI

- Data/Model parallel supported(3D)
- 3D Tensor parallelism
- ZeRO-Offload: model data supported
- MoE specific parallelism
- RLHF Demo: ColossalChat

End of Parallel strategies

Result & Evaluations

- NLU
 - SuperCLUE
- NLG
- NLI

Finetuning

- Target and issues
- Instruct finetuning
- Finetuning for specific task
- Parameter efficient finetuning

Target and issues

- Target
 - From pattern completion to real world tasks
- Issues
 - Instruction following
 - Hallucination
 - Toxicity and ethics
 - Securities

Instruct Finetuning

Key to success

- Number of finetuning datasets:
 - scaling from 62 text datasets to 18K
- Model scale: 137B LaMDA-PT
- Natural language instructions

Finetune on many tasks (“instruction-tuning”)

Input (Commonsense Reasoning)

Here is a goal: Get a cool sleep on summer days.
How would you accomplish this goal?
OPTIONS:
-Keep stack of pillow cases in fridge.
-Keep stack of pillow cases in oven.

Target

keep stack of pillow cases in fridge

Input (Translation)

Translate this sentence to Spanish:
The new office building was built in less than three months.

Target

El nuevo edificio de oficinas se construyó en tres meses.

Sentiment analysis tasks

Coreference resolution tasks

...


Inference on unseen task type

Input (Natural Language Inference)

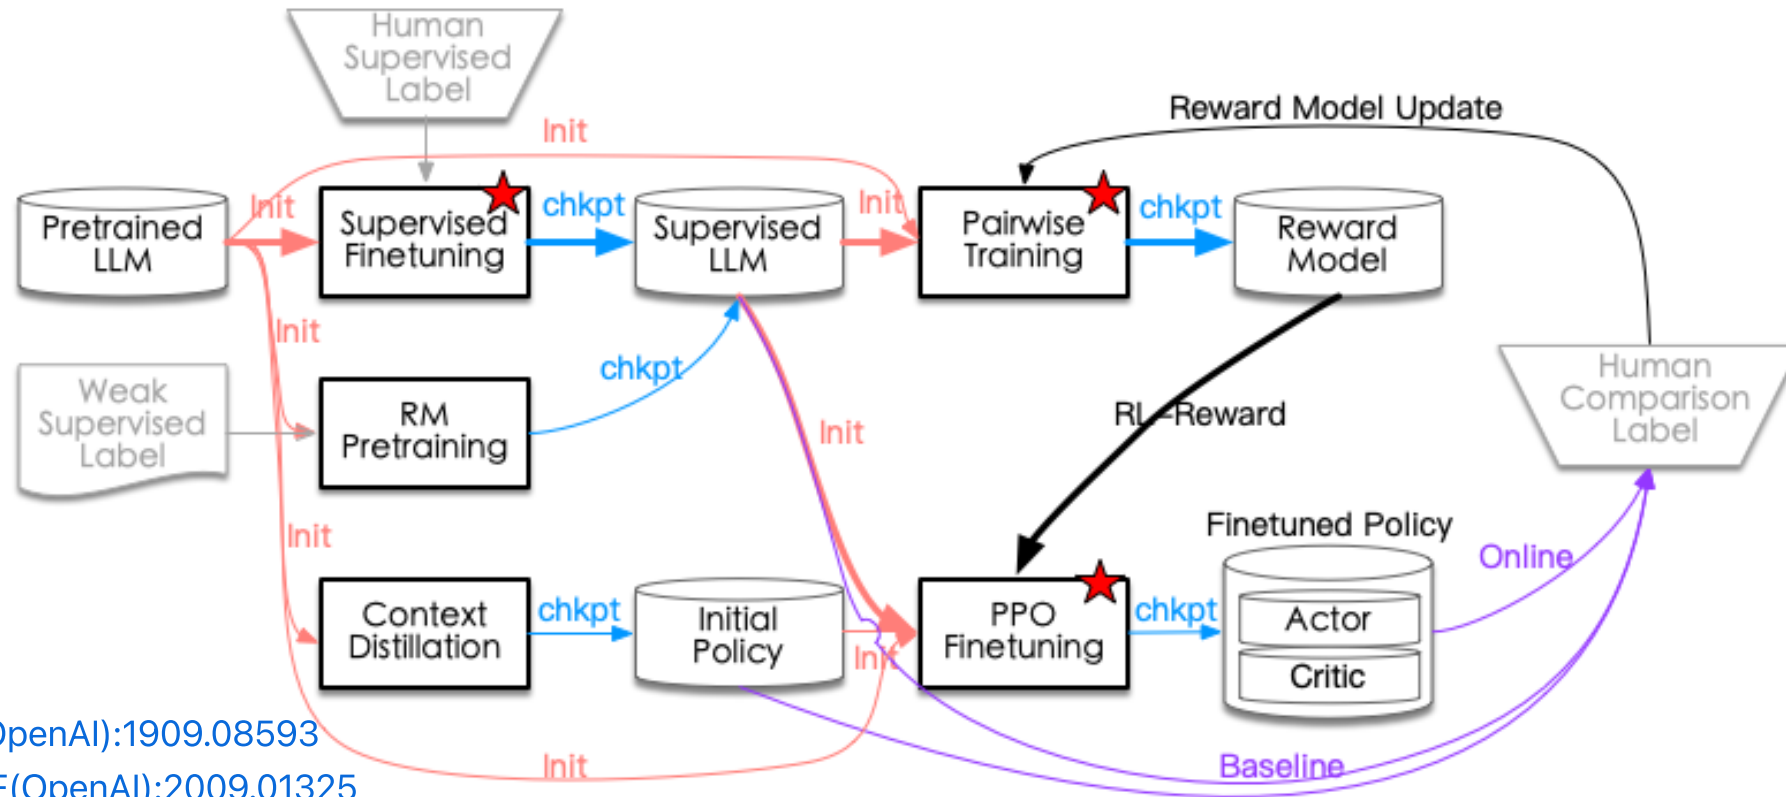
Premise: At my age you will probably have learnt one lesson.
Hypothesis: It's not certain how many lessons you'll learn by your thirties.
Does the premise entail the hypothesis?
OPTIONS:
-yes -it is not possible to tell -no

FLAN Response

It is not possible to tell



Finetuning for specific task



Finetune from HF(OpenAI):1909.08593
Summarize from HF(OpenAI):2009.01325
HHHA(Anthropic):2112.00861
InstructGPT(OpenAI):2203.02155
HHA(Anthropic):2204.05862
Sparrow(Deepmind):2209.14375

Problems from Supervised Finetuning(SFT)

- Learning only the task format and the way to response for the format
- Knowledge labeled but not in the LLM leads to more hallucination
- Knowledge in the LLM but labeled as `don't know` leads to withhold information

What we want from finetuning:

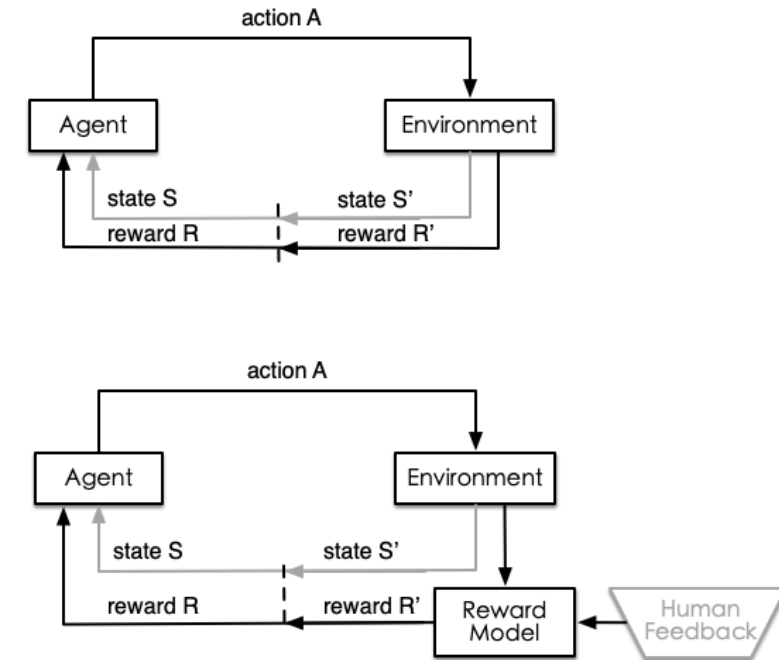
outputs its(LLM's) state of knowledge with the correct amount of hedging and expressing its uncertainty

Advance of RL to SFT for truthfulness

- LLMs know what they know
 - Calibrated probability, uncertainty
- RLHF can leverage the self-awareness
 - Design reward function: correct answer=1 , don't know=0 and wrong answer=-4
 - RL learn optimal threshold of probability to maximize the reward
- No oracle for the correctness, delegate to Reward Model
 - Reward model: **relative criteria** trained by pairwise loss from human feedback
 - Open problem: true probabilities of everything?
 - Open problem: go beyond things that labelers can easily do
 - Verification is easier than generation

More on Reinforcement Learning


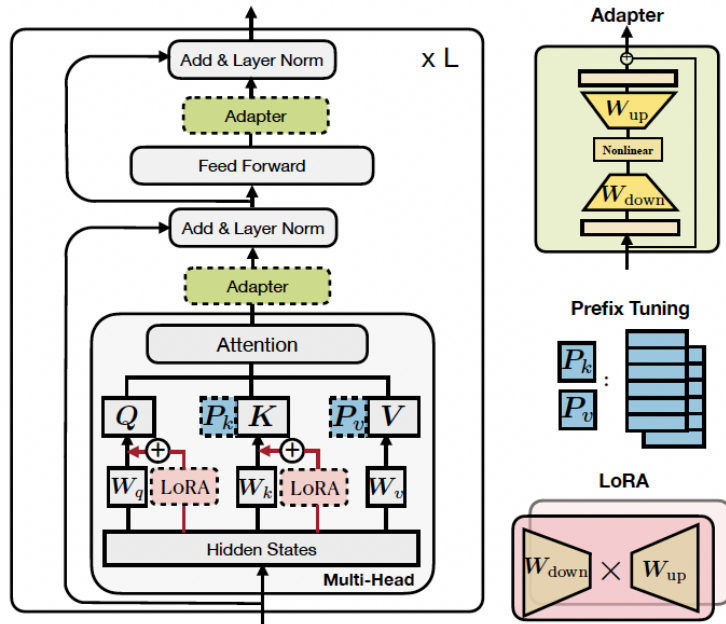
- Catalog of algorithms(**PPO belongs**)
 - World model or **model free**
 - Value-based or **policy-based(actor-critic)**
 - MC or **TD bootstrapping**
 - Off-policy or **on-policy**
 - Deterministic or **stochastic** policy
- Design consideration
 - **Sample efficiency**: Off-policy > On-policy
 - **Stability & Convergence**: Deadly Triad issue
 - **Explore & Exploit**: Random at episode beginning
 - **Bias & Variance**: Advantage Function



Parameter Efficient Finetuning

- Design aspects
- Primary implementations
 - Adapter
 - Prefix-tuning
 - LoRA

Methods and performance on Summerrization task




PEFT illustration and performance comparison (Source: [Junxian He, et.al](#))

Design aspects

- **Finetuned Modules:**
 - Attention-key/value matrix: LoRA(Q/V)
 - Attention-head: Prefix-Tuning(K/V)
 - Attention: Adapter
 - After FFN: Adapter
- **Other aspects:**
 - Multi-task consideration
 - Task related head

Adapter

- Implementation & training notes
 - $h \leftarrow h + f(hW_{\text{down}})W_{\text{up}}$
 - Parameter scale:
 $2 \times L \times (r \times d + r + d), r \ll d$
 - Adapt after FFN sub-layer works too
- Results
 - Finetune BERT for 26 classification Tasks
 - 3.6% parameters for 0.4% GLUE performance gap
 - Ablation: fewer layers adapted -> worser performance



Prefix-Tuning

- Implementation
 - $\text{head} = \text{Attn}(XW_Q, [P_K; XW_K], [P_V; XW_V])$
 - Reparameterization for finetuning stability:
 - $[P_K, P_V] = \text{MLP}_\theta(P^E)$
 - P^E : prefix embedding
- Parameter scale:
 - Vanilla: $|P| \times d \times 2 \times L$
 - Reparameteration: $|P| \times d + d \times H + H \times d \times 2 \times L$

More on Prefix-Tuning: Training and scaling

- Training
 - Initialization:
 - Real/high frequency words activation
 - *Task relevant words / Classification labels*
 - LM Head: Next Token/Class Label
- Results & discussion
 - Finetuning 0.1% \sim 3% parameters, comparable or better performance
 - Optimal prefix length varies: longer for more complex tasks
 - Reparameterization works task-dependently

PrefixTuning: Optimizing Continuous Prompts for Generation, Stanford, 2021

PromptTuning: The Power of Scale for Parameter-Efficient Prompt Tuning, Google, 2021

P-Tuning v2: Prompt Tuning Can Be Comparable to Fine-tuning Universally Across Scales and Tasks, Tsinghua, 2021

LoRA

- Implementation & training notes
 - Transformer: $W = W_0 + (BA)^T, A \in R^{r \times d}, B \in R^{d_h \times r}, r \ll \min\{d_h, d\}$
 - W_Q and W_V considered, parameter scale: $2 \times 2 \times d \times r \times L$
 - Modularized: Embedding , Linear , MergedLinear , Conv2D
 - Initialization: A kaiming-random, B zeros
 - Weight merged for inference efficiency
- Results
 - For 175B GPT-3 finetuning: 0.01% parameters , on par or better results
 - No additional inference computation and latency
 - Additivity for **finetuning merge** and **incremental update**

More on LoRA: which part to update & rank settings

	# of Trainable Parameters = 18M						
Weight Type Rank r	W_q 8	W_k 8	W_v 8	W_o 8	W_q, W_k 4	W_q, W_v 4	W_q, W_k, W_v, W_o 2
WikiSQL ($\pm 0.5\%$)	70.4	70.0	73.0	73.2	71.4	73.7	73.7
MultiNLI ($\pm 0.1\%$)	91.0	90.8	91.0	91.3	91.3	91.3	91.7

	Weight Type	$r = 1$	$r = 2$	$r = 4$	$r = 8$	$r = 64$
WikiSQL($\pm 0.5\%$)	W_q	68.8	69.6	70.5	70.4	70.0
	W_q, W_v	73.4	73.3	73.7	73.8	73.5
	W_q, W_k, W_v, W_o	74.1	73.7	74.0	74.0	73.9
MultiNLI ($\pm 0.1\%$)	W_q	90.7	90.9	91.1	90.7	90.7
	W_q, W_v	91.3	91.4	91.3	91.6	91.4
	W_q, W_k, W_v, W_o	91.2	91.7	91.7	91.5	91.4

End of Parameter Efficient Finetuning

Steering the decoding process of LLM

- Decoding strategies
- Prompt engineering

Decoding strategies

- Temperature in decoding: $p(w_i|w_{<i}) = \frac{\exp(o_i/T)}{\sum_j \exp(o_j/T)}$, o_i logits from LLM
- Maximal Likelihood Search
 - Greedy search: $w_i = \arg \max p(w_i|w_{<i})$, eq. to $T = 0$
 - Beam search: $w_i \in \text{TopN } p(w_i|w_{i-1}, w_{<i-1})p(w_{i-1}|w_{<i-1})$
- Sampling
 - top-K sampling: $w_i = \text{sample TopK } p(w_i|w_{<i})$
 - top-p(Nucleus) sampling: $w_i = \text{sample TopK}_i \sum_{w_i < K_i} p(w_i|w_{<i}) \geq p$
 - Repetition penalized sampling
- Guided decoding
 - $\text{score}(x_{t+1}, b_t) = \text{score}(b_t) + \log p(x_{t+1}) + \sum_i \alpha_i f_i(x_{t+1})$

Prompt engineering

- One-shot interaction
 - Instruction/Zero-shot Prompt
 - Few-shot Prompt
 - In-context Learning(Prompt)
 - Chain-of-Thought
- Recursive interaction
 - MRKL: [Thought/Action/Action Input/Observation]+ , zero-shot, access tools
 - Self-ask: Followup question? [Question/Answer]+ Final Answer , few-shot
 - ReACT: [Thought/Action/Observation]+ , few-shot, access tools

More on In-context Learning

What matters?

- Examples template(instruct/CoT) & order: yes
- Label space & Input distribution(diversity): yes
- Exact {Question, Answer} pair: no/yes

Why it works?

- Interpretation from Topic Model: $P(o|p) = \int_z P(o|z, p)P(z|p)dz$
- Induction head: [a] [b] ... [a] -> [b]

[A Mathematical Framework for Transformer Circuits, 2021, Anthropic](#)

[How does in-context learning work?, 2022, Stanford](#)

[Rethinking the Role of Demonstrations: What Makes In-Context Learning Work?, 2022, Meta](#)

[Larger language models do in-context learning differently, 2023, Google](#)

Augmentation and Plugins

- Augmented Language Models
- Automatic prompting
- Plugins

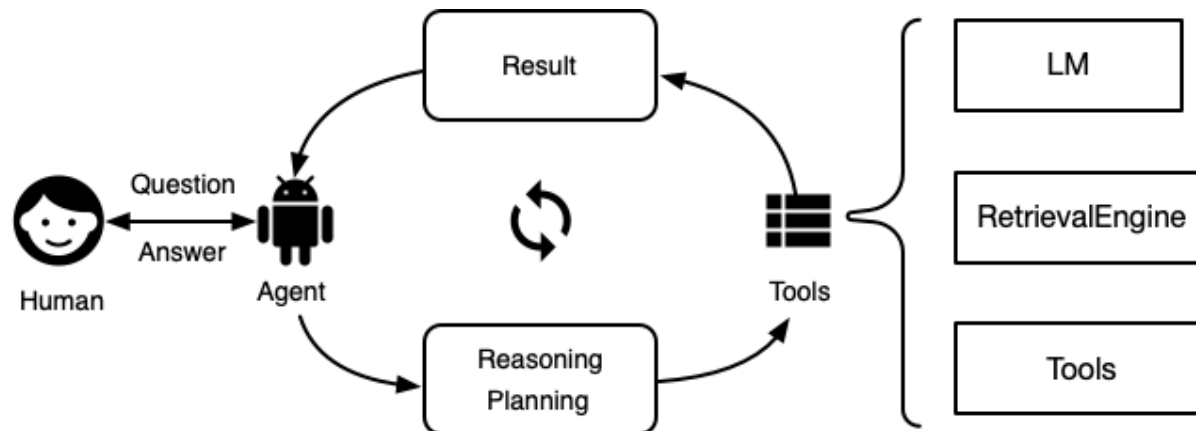
Augmented Language Models

- Problems in vanilla LLMs
- Augment aspects
 - Reasoning/Planning
 - Tool usage

Problems in vanilla LLMs

- Compression of world knowledge
- Context limitation
- No up-to-date knowledge
- Hallucinations

Augmented Language Model Overview



- 2 Steps
 - What: Planning/Reasoning
 - How: Tool usage
- 3 Teaching Methods
 - In-context prompt:
 - One-shot/Recursive
 - Finetuning
 - Reinforcement Learning:
 - Hardcode/Human Feedback

Planning by self-talk: Self-ask

- Multi-hop question and compositionality gap
- Self-ask prompt:
 - Few-shot prompt scaffold:


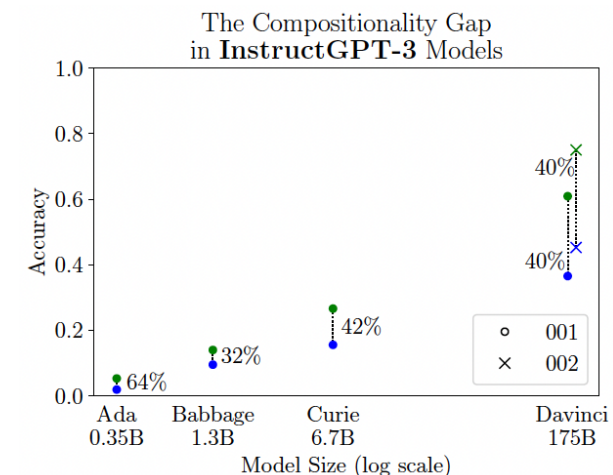
```
Question: {{Question}}  
Are follow up questions needed here: Yes  
Follow up: {{Follow-up-question}}  
Intermediate answer: {{Intermediate-answer}}  
So the final answer is: {{Final-answer}}
```

```
Question: {{Question}}  
Are follow up questions needed here: {{LLM-gen}}
```

- Few-shot prompt with search engine:
 - The same prompt as before
 - Generated by query search engine:

```
Intermediate answer: {{Intermediate-answer}}
```

Accuracy on *Compositional Celebrities* 2-hop questions



Planning by self-talk: Self-ask prompt

Question: Who lived longer, Theodor Haecker or Harry Vaughan Watkins?

Are follow up questions needed here: Yes.

Follow up: How old was Theodor Haecker when he died?

Intermediate answer: Theodor Haecker was 65 years old when he died.

Follow up: How old was Harry Vaughan Watkins when he died?

Intermediate answer: Harry Vaughan Watkins was 69 years old when he died.

So the final answer is: Harry Vaughan Watkins

Question: Who was president of the U.S. when superconductivity was discovered?

Are follow up questions needed here: Yes.

Follow up: When was superconductivity discovered?

Intermediate answer: Superconductivity was discovered in 1911.

Follow up: Who was president of the U.S. in 1911?

Intermediate answer: William Howard Taft.

So the final answer is: William Howard Taft.

Planning by self-talk: Self-ask prompt with search engine

Question: Who lived longer, Theodor Haecker or Harry Vaughan Watkins?

Are follow up questions needed here: Yes.

Follow up: How old was Theodor Haecker when he died?

Intermediate answer: Theodor Haecker was 65 years old when he died.

Follow up: How old was Harry Vaughan Watkins when he died?

Intermediate answer: Harry Vaughan Watkins was 69 years old when he died.

So the final answer is: Harry Vaughan Watkins

Question: Who was president of the U.S. when superconductivity was discovered?

Are follow up questions needed here: Yes.

Follow up: When was superconductivity discovered?

Intermediate answer: Superconductivity was discovered in 1911.

Follow up: Who was president of the U.S. in 1911?

Intermediate answer: William Howard Taft.

So the final answer is: William Howard Taft.

Conditional LM on retrieved documents: REALM

Bootstrapping with self-supervision for reasoning: StAR

Learning to code: Codex

Teach LM to use search engine: WebGPT

End of Augmented Language Models

Plugins

- Plugins ecosystem
- Primary implementations
 - Langchain Agent/Tool
 - ChatGPT Plugins
 - Fixie
 - Other paradigm proposal

Plugins ecosystem

- Tool as a service
- From SEO to LMO
- Orchestration by LLM

Langchain Agent/Tool

ChatGPT Plugins

Fixie

Other paradigm proposal

End of Plugins

Thanks & QA?